


Triangle Listing

Team 02

The Problem: List All Triangles in an Undirected Graph

- Undirected simple graph
- A triangle: three nodes with pairwise edges
 - $\{A, B, C\}$, $\{A, B, D\}$, $\{A, D, E\}$
- Listing vs counting
- Useful for real-world network analysis



Overview

Algorithm Implementations

Instrumentation & Benchmarking

Graph Generation and “Adjusting” Real World Graphs

Optimization Ideas

Edge Iterator: Finds Shared Neighbors of Edges

- For every edge $\{u, w\}$
- If node v is in both $\text{Adj}(u)$ and $\text{Adj}(w)$
- Then $\{u, v, w\}$ is a triangle

- We store the graph as adjacency lists
- Triangle test: $\text{Adj}(u) \cap \text{Adj}(w)$
 - Sorted $\text{Adj}(u)$ and $\text{Adj}(w)$ can be intersected in linear time
 - We use quicksort to preprocess them
 - Included in the total execution time

- Only consider edges (u, w) s.t. $u < w$, to avoid counting duplicate triangles
- Equivalent to transforming the undirected graph to a directed one

Forward: Dynamically Grow Neighbor Lists

- Explicitly store in-edges of w in a dynamic structure $A(w)$
 - After (u,w) is visited, add u to $A(w)$
 - $|A(w)| \leq d_{\text{in}}(w) \leq |\text{Adj}(w)| = d(w)$
 - $A(w)$ uses the same data structure as $\text{Adj}(w)$
 - $\text{Adj}(w)$ is sorted $\Rightarrow A(w)$ is also sorted
- Triangle test: $A(u) \cap A(w)$ instead of $\text{Adj}(u) \cap \text{Adj}(w)$

Algorithm 1: *forward*

Input: ordered list of vertices $(1, \dots, n)$, Adjacencies $\text{Adj}(v)$

Data: Node Data: $A(v)$;

for $v \in V$ **do**

$A(v) \leftarrow \emptyset$

for $s \in (1, \dots, n)$ **do**

for $t \in \text{Adj}(s)$ **do**

if $s < t$ **then**

foreach $v \in A(s) \cap A(t)$ **do**

 output triangle $\{v, s, t\}$;

$A(t) \leftarrow A(t) \cup \{s\}$;

Forward Hashed: Test Neighbors with Hash Tables

Forward Hashed

- Use a hash container for $A(w)$
- Triangle test: $A(u) \cap A(w)$
 - Use the smaller hash table to probe the larger one
 - If hash table lookups take $O(1)$ time, the intersection takes $O(\min\{d_{in}(u), d_{in}(w)\})$

Instrumentation & Benchmarking

- Templated the Index type used in the implementations to count Integer operations
- Benchmarking framework:
 - Loads the graph and create helper data structures.
 - Run Instrumented version
 - Reload the graph
 - Do \$WARMUP warmup runs
 - Run \$PHASES phases with each \$RUNS iterations and time the cycles

```
./benchmark -num_warmups $WARMUP -num_runs $RUNS -num_phases $PHASES -o $OUTPUTDIR/$graph.csv  
-algorithm edge_iterator,forward,forward_hashed -graph $INPUTDIR/$graph.txt
```

Benchmarking - Hardware

- Intel Core i9-9900K - Coffe Lake
- 4 ALU-Ports: Peak Performance 4 Integer Arithmetic ops per cycle (4x4 Vec)
- Cache Size:
 - L1: 32 KiB, 8-way set associative per core
 - L2: 32 KiB, 8-way set associative per core
 - L3: 16 Megabyte, 16-way set associative shared

Real-World Graphs


- German Road Network (2x larger)
- Computer Science Co-Authorship (roughly equal)
- Google Contest ((2x, 9x) larger)
- HEP (High Energy Particle Physics) Citation graph (roughly equal)
- Web page links nd.edu domain (equal)
- US Patent Citations (equal)

Graph Transformation and Filtering

- Transformations
 - Output a list of edges for each node
 - Directed -> Undirected
 - Treat edges as undirected and add missing endpoints to respective adjacency lists
- Filtering
 - Remove Multi - Edges
 - Remove Loops


cycles

Runtime



ops

Op Count

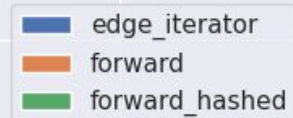


ops/cycle

Performance

10^0

10^{-1}




HEP_LITERATURE 27770,352285
COMP_SCIENCE_AUTHORS 317080,1049866
WWW_NOTRE_DAME 325729,1090108
GOOGLE_CONTEST 875713,4322051
US_PATENTS 3774768,16518947
GERMAN_ROAD_NETWORK 11548845,12369181

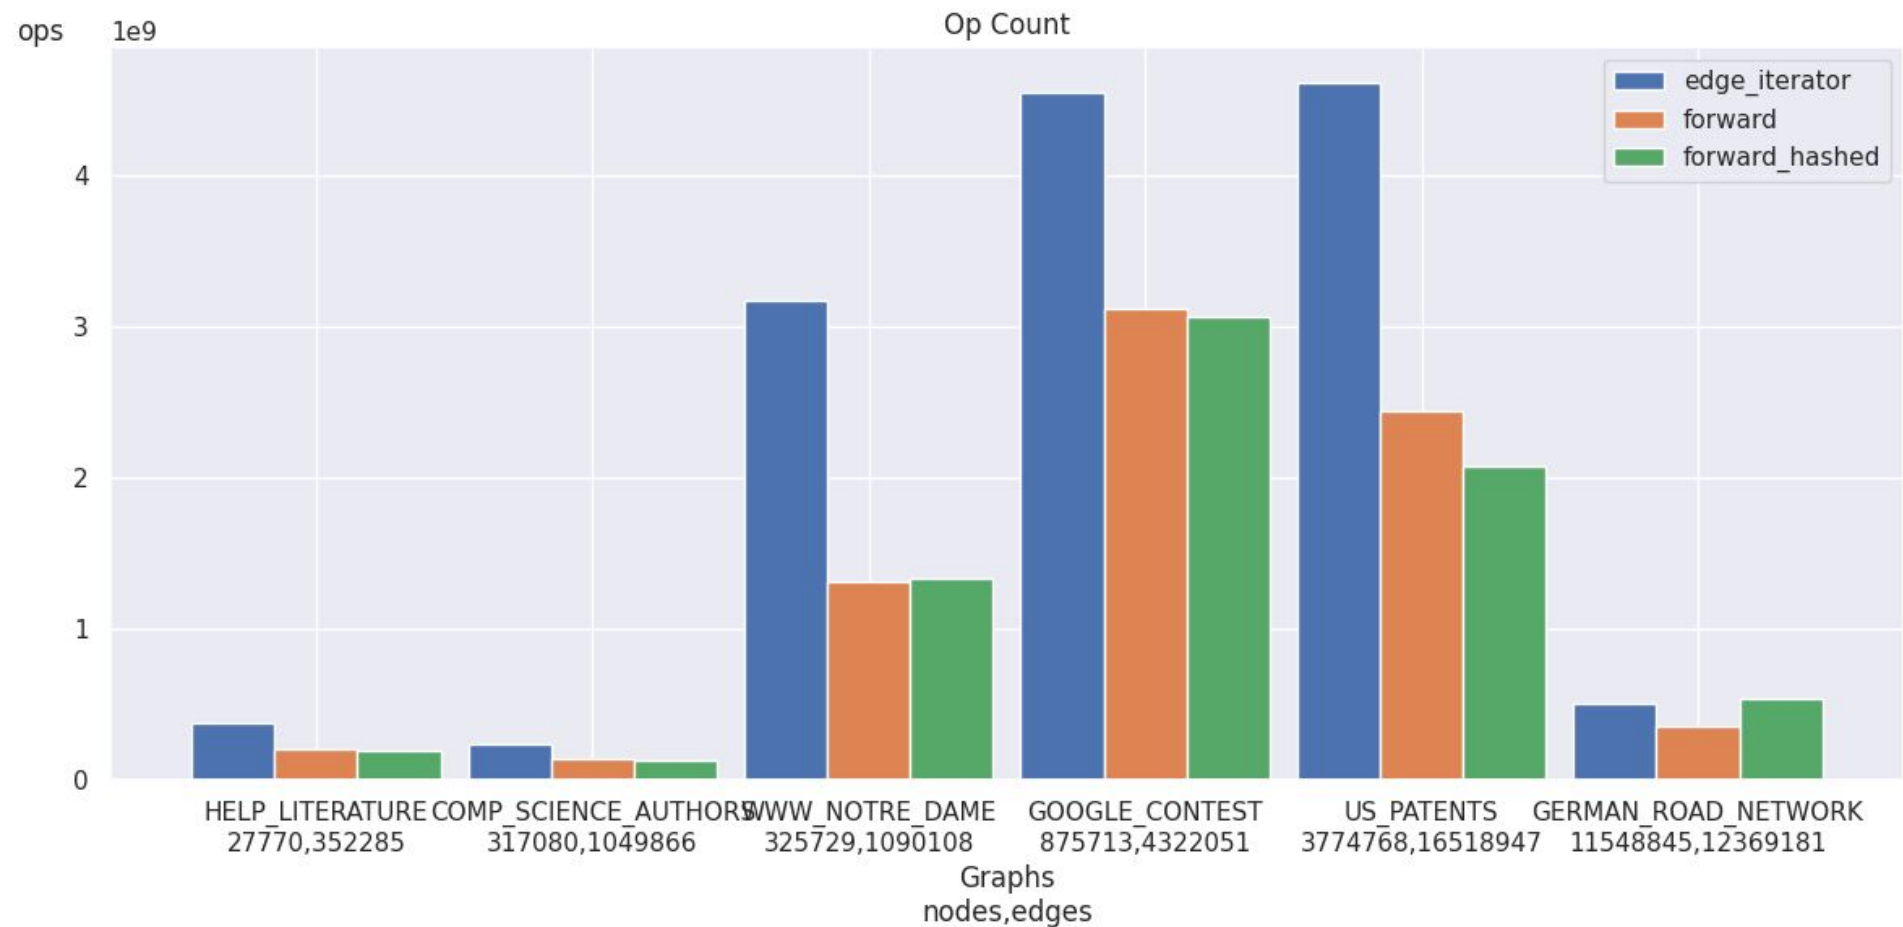
Graphs
nodes,edges



cycles


Runtime





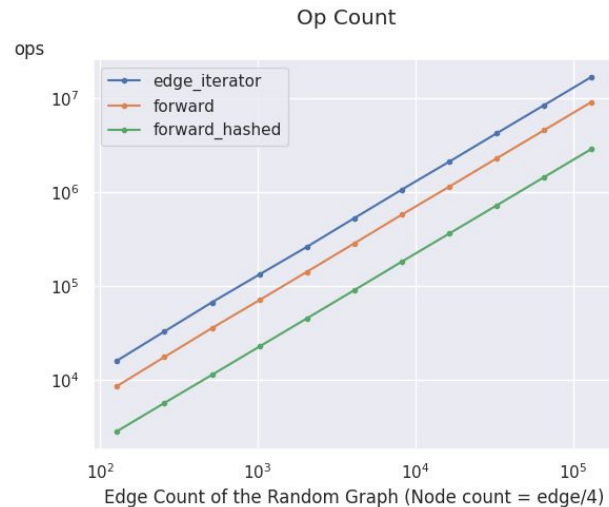
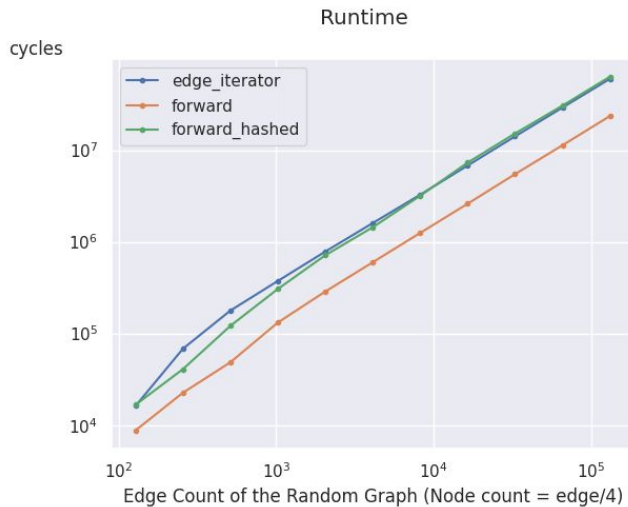
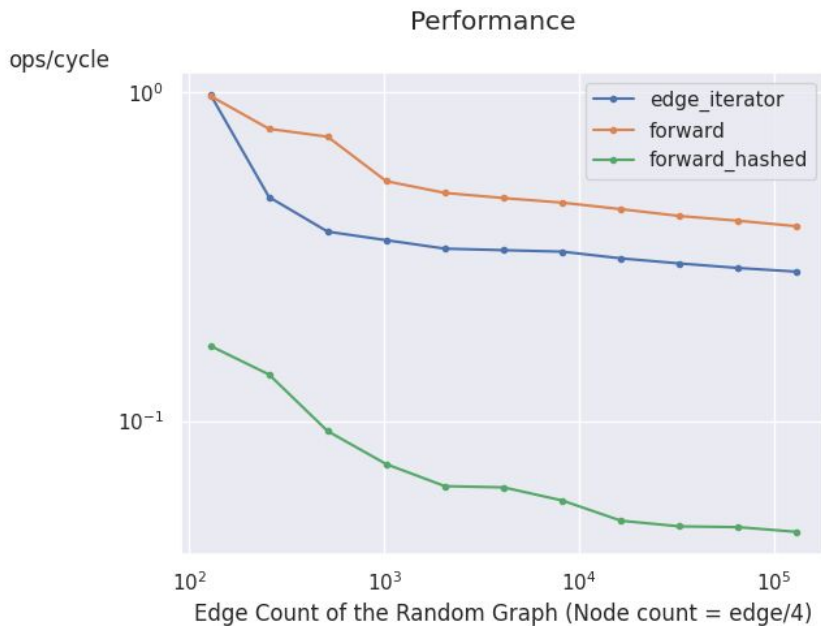
ops/cycle

Performance



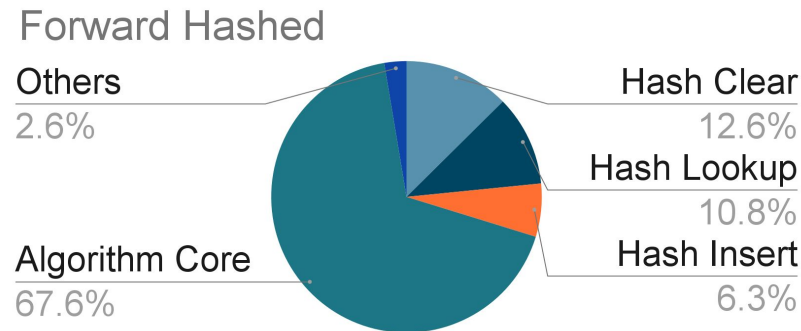


Random Graph Results

- Randomly distribute connect nodes until all edges have been created



Profiling Result

- Algorithm cores take most of the time.
- There is also space for optimization in quick sort and hash table.



Optimization Ideas

- Integer size
 - 8-byte ints or smaller ints
- Indirections
 - Pointers in structs and linked lists
 - Reduce the number of indirections in the data structures
- Data Structures
 - Data structure optimizations towards different graph properties (density, max degree, etc.)
 - More compact
 - More efficient for preprocessing, triangle tests and hash table operations

Optimization Ideas

- Function Inlining
- ILP
 - 4 ports each with 1 ALU execution unit
 - Set intersections are independent, but the naive implementation of set intersection is blocking
 - Loop unrolling
- Cache
 - Temporal locality of adjacency lists: blocking
 - Update dynamic structure
- Branching
 - Branching: quicksort, undirected→directed transformation, set intersection, hash table lookup
 - Reduce branch mispredictions

Optimization Ideas

- Hash Table
 - Hash functions: modulo
 - Hash container size
 - Collision resolution: separate chaining vs. open addressing
- SIMD
 - Use SIMD instructions to accelerate set intersections and hash table operations

Questions

- What exactly should be counted as an operation for the performance plots?
- Benchmark listing vs counting?
- Warmup runs?
- Should we time resetting the helper datastructure?
- How should we un-sort the edges again, this basically doesn't allow us to have more than one run per phase (which is probably fine for big graphs).
- Memory allocation?
- To what extent can we change the graph representation and operations so that the algorithms can still be considered as the algorithms given in the paper?
- C++ for template in the main algo body?